



Real-Time Speech Translation Application Feasibility

Overview: Crawl, Walk, Run Phases of the App

This project envisions an **audio speech translator** that evolves in stages. In the **“crawl” phase**, the focus is on offline audio file translation: e.g. taking a recorded **MP3/WAV in Russian** and outputting an **audio file in English** (or vice versa). The next **“walk” phase** would enable real-time translation locally – for example, a script or app that listens to a live **Zoom call** and produces a translated audio stream on the fly. Finally, the **“run” phase** targets a full **web application** where a user can submit a **video/stream link** (YouTube, Zoom meeting, etc.) and receive a real-time translated audio playback (dubbing) alongside the original video. A **slight delay ($\approx 1-2$ seconds)** is acceptable in all phases to allow processing latency. This staged approach ensures we start simple (file translation) and incrementally tackle the harder real-time streaming scenarios.

Key Technology Components and Feasibility

Building a real-time translation app is **technically feasible** today, thanks to advances in speech and translation APIs. The solution breaks down into a pipeline of three core components: **Speech Recognition**, **Machine Translation**, and **Speech Synthesis**. Modern APIs and frameworks make it possible to chain these steps with low latency, or even perform them end-to-end. Below is a summary of each component and available tools:

- **Speech-to-Text (ASR)** – This converts spoken Russian or English audio into text. Industry APIs like **Google Cloud Speech-to-Text**, **Microsoft Azure Speech**, and **Amazon Transcribe** offer high accuracy for both English and Russian and can operate in real-time streaming mode. For example, Amazon Transcribe can stream live audio and output text transcripts (it even supports automatic language detection) ¹ ². Open-source models are also an option; **OpenAI’s Whisper** (2022) is a state-of-the-art multilingual ASR model that can transcribe speech in ~100 languages and even directly **translate foreign speech to English text** ³. Whisper could handle Russian-to-English transcription+translation locally, though it would still require a separate step to go *into* Russian (since Whisper’s translation mode is primarily to English). Overall, ASR for English and Russian is a solved problem – these languages are well-supported by existing speech recognition models, with near human-level accuracy in favorable conditions.
- **Machine Translation (MT)** – This translates between Russian and English text. All major cloud providers have neural translation engines that support **Russian↔English** (among dozens of other languages). **Google Translate** and **Microsoft Translator** are highly developed for this language pair, offering fluent translations via APIs. **Amazon Translate** supports 75+ languages (as of mid-2022) with continually improving quality ⁴. These services can be used on-the-fly in the app to translate recognized text with low latency (typically sub-second for sentence-length inputs). For instance, Microsoft’s cloud Translator or Azure Cognitive Services can translate streaming transcripts in real time. If needed, open-source MT models (like MarianNMT or those on HuggingFace) could be

employed, but leveraging cloud APIs ensures up-to-date quality and avoids heavy model training. Given the maturity of **English-Russian** machine translation (it's a common high-resource language pair), the translations should be reasonably accurate for general domains.

- **Text-to-Speech (TTS)** – This synthesizes natural-sounding speech in the target language (Russian or English) from the translated text. Again, cloud APIs excel here: **Amazon Polly**, **Google Cloud Text-to-Speech**, and **Azure Speech** all offer lifelike voices in many languages. For example, Amazon Polly has multiple English voices and supports Russian (“Русский”) voice options as well ⁵. These voices can be used to **speak out the translated text in real time**. Cloud TTS can return audio quickly (often generating speech faster than real-time). Azure and Google likewise have high-quality neural voices for both English and Russian. If an on-device solution is needed, there are open libraries (e.g. Coqui TTS or eSpeak with higher-quality voice models), though quality may be lower. The ideal is to use a cloud TTS for natural prosody. It's worth noting that Azure's Speech service even allows **voice customization** (training a custom voice or using **neural voice cloning** with permission), which could theoretically be used to mimic a speaker's voice in another language ⁶ ⁷, though that's an advanced optional feature.

Using these components, we can construct a **speech-to-speech translation pipeline**. The **technical feasibility is well-demonstrated** by existing solutions. For instance, **Amazon Web Services (AWS)** published a 2024 demo using **Transcribe + Translate + Polly** in a Python app to produce a “*near-real-time speech-to-speech translator*”, requiring “*zero ML experience*” ⁸. They describe streaming audio into Transcribe, sending each transcribed segment to Amazon Translate, and then passing the result to Polly to play the translated voice output immediately ⁹ ¹⁰. This kind of pipeline can achieve only about ~1–2 seconds of lag, primarily due to waiting for a pause or end-of-sentence to finalize the translation and the TTS synthesis time. A small buffering delay is acceptable as noted, and indeed AWS calls it “near real-time.” **Microsoft Azure** offers an even more integrated route: their **Speech Translation API** can directly take audio input and output *translated audio* in a target language using neural voices, all in one streaming call ¹¹. This means Azure's service internally does the ASR, translation, and TTS for you. It was recently enhanced to support *multi-language input* (auto-detect speakers' languages) and continuous translation in one session ¹² ¹³ – ideal for meetings with mixed languages. **Google** isn't far behind; they provide a **Media Translation API** for streaming speech translation. Google's API similarly promises “*low-latency streaming translation*” combining their speech recognition and Translate technologies ¹⁴ ¹⁵. In short, all the building blocks are available via robust APIs. By using these services (instead of training our own models from scratch), a custom app can be developed relatively quickly. The heavy lifting (speech recognition accuracy, linguistic translation, voice quality) is handled by these cloud AI models which are continuously improving. The developer mainly needs to orchestrate the calls and handle audio input/output.

Latency Considerations: The goal of ~1–2 seconds delay is quite achievable. Streaming ASR APIs output interim results within fractions of a second after a phrase is spoken. The translation of one sentence worth of text typically takes well under a second on cloud servers. TTS can synthesize a sentence of audio in perhaps 200–500 milliseconds. The key is to process audio in **chunks** (e.g. sentence or phrase segments) rather than waiting for an entire speech. Real-time translation systems usually work on a rolling basis: as soon as you speak a clause and pause, the translation is spoken out. A brief delay is inherent, but users generally tolerate a couple seconds in live conversation (similar to a human interpreter's lag). In fact, Microsoft's Skype Translator (one of the early real-time translators) introduced about a 1–2 second delay and was still considered effective in enabling conversation. Today's models are faster and can often start

speaking the translation even *before* the speaker has finished a long sentence (by using smart buffering/prediction or simply breaking on pauses). Therefore, meeting the 1–2 second lag requirement is feasible with careful streaming pipeline implementation.

Development Approach: We will indeed *develop our own app UI/logic* but **leverage these API services** for the core translation intelligence (this is the sensible approach – building these models from scratch would be an enormous task). The app can be structured such that it streams audio to the cloud (or runs local models if offline mode is needed) and then plays back the translated audio to the user. Modern frameworks (Python with async I/O as shown by AWS's example, or Node/Python using Google's streaming gRPC, etc.) can handle audio streaming. Also, some platforms provide SDKs – e.g. Azure's SDK or Google's client libraries – to simplify real-time use.

Figure: High-level architecture of a speech-to-speech translation pipeline (from an AWS example). Audio from the user is transcribed by a Speech-to-Text engine, translated to the target language text, then sent to a Text-to-Speech engine to produce translated audio output ¹⁶.

Use Case 1: Translating YouTube Videos (Near Real-Time Dubbing)

Scenario: For a YouTube video (either a recorded video or a live stream), the user wants to hear the audio in another language (Russian ↔ English) while still seeing the original video picture. This essentially means **dubbing the video** with an AI-generated translated voice.

Feasibility: This is quite feasible with today's tech. In fact, companies and researchers are actively doing this for content localization. For instance, Google's Area 120 project "Aloud" is an AI dubbing tool for YouTube creators (it transcribes, translates, and narrates videos into other languages) ¹⁷. YouTube has even started supporting multiple audio tracks on videos, and auto-dubbing features are emerging ¹⁸ ¹⁹. Our approach would be: 1. **Extract or capture the audio** from the YouTube video. If it's a recorded video, the app can download the audio stream (using YouTube's API or tools) quickly. For a live stream, the app can capture the live audio via the streaming URL.

2. **Translate the audio to a new audio track** using the speech pipeline: The audio is fed through ASR → MT → TTS as discussed. This can be done in streaming fashion for live video (so the translated voice plays almost concurrently with the video, a second or two delayed). For an existing video file, one could also process it faster-than-real-time (since it's not truly live) by chunking the audio and translating, then output the full translated track.

3. **Sync the translated audio with the video.** In a real-time scenario, the app would play the video for the user but mute the original audio and play the synthesized translated audio instead. (Alternatively, it could mix at lower volume if the user wants to hear the original lightly, though usually for dubbing we replace it.) If processing offline, one could output a new video file with the dubbed audio track.

Technical aspects: A slight delay is acceptable; for a live stream the app might buffer ~2 seconds of video to ensure the translated audio can keep up without gaps. The translation doesn't need to be perfectly word-for-word with the original; we aim for semantically equivalent output delivered smoothly. Since 1-2 second delay is fine, we don't have to aggressively cut off sentences – we can wait for punctuation or pauses which improves translation quality. Modern ASR like Whisper or Google's will produce punctuated transcripts, which helps in determining when to translate and speak. Also, many YouTube videos speak in sentences/phrases that allow short pauses.

For **quality**, the output voice can be chosen to suit the content – e.g. a formal voice for a documentary, or a more upbeat voice for entertainment. The TTS voices allow customization (volume, rate, etc.) via SSML if needed ²⁰. One challenge is **emotional tone or speaker identification** – a single audio track doesn't convey different speakers distinctly (unless we intentionally use different voices for different detected speakers, which is complex and often not necessary for understanding). In most cases, a single translated voice is acceptable. If the YouTube video has multiple speakers and it's important (say a dialogue), one could attempt to detect speaker changes and alternate voices (male/female, etc.) in the synthesis stage. This is an advanced feature and not essential for an MVP, but it's technically possible with speaker diarization from the ASR and using different voice IDs accordingly.

Existing solutions & demand: There is strong demand for translating video content. Creators want to reach wider audiences, and viewers want to consume content in their own language. Aside from Google's efforts, startups like **Papercup** offer AI dubbing services for videos, focusing on translating YouTube or media content into multiple languages with synthetic voices. This indicates a market need – many companies are using AI to localize e-learning videos, marketing content, etc., though typically not in real-time. Our edge would be doing it *on-the-fly* for any given video link. Tools like **Notta.ai** and some browser extensions already let users auto-generate translated subtitles or even dub audio for YouTube, which validates the concept. Overall, from a technical viewpoint, processing a YouTube video is actually easier than a live conversation because the speech is one-way and not interactive. The translation pipeline just needs to keep up with playback. Therefore, **implementing YouTube video translation first makes sense** – it allows testing the core translation quality and latency in a controlled setting.

Use Case 2: Live Translation for Zoom Meetings

Scenario: In a Zoom meeting (e.g. an English speaker and a Russian speaker in a call), our app would provide **real-time translated audio** so that each participant can understand the other in their preferred language. Essentially, it functions like an automated interpreter for Zoom. The output could be delivered as an audio feed that the user hears while watching the meeting.

Feasibility & Integration: Technically, this is more complex than the YouTube case because it's **bi-directional** and interactive. However, it's feasible and in fact already starting to appear in the market. Zoom itself has introduced **translated live captions** (an add-on service that shows real-time subtitles translated into 30+ languages) ²¹ ²². Zoom's built-in solution currently focuses on captions (text), not voice dubbing. This leaves room for an app that **speaks** the translation aloud. One approach is a companion app that the user runs alongside Zoom: the app listens to the meeting's audio and whenever someone else speaks Russian, it outputs English audio for our user (and vice versa).

To implement this, the app needs access to the Zoom audio stream. There are a few ways: - **Via System Audio:** If the user is on a desktop, the app can capture the system's sound output (the Zoom meeting audio) through an audio driver or loopback. Simultaneously, it can capture the user's microphone input if we also translate the user's speech for others. The app would then process each direction through the pipeline. - **Zoom's API/SDK:** Zoom provides an SDK for developers (and a concept of "Zoom Apps"). One could use the Zoom SDK to join the meeting as a bot participant that has access to raw audio. However, Zoom's SDK may have limitations on programmatically extracting audio in real-time due to security. Alternatively, Zoom offers a feature called "Live Interpretation" where human interpreters join on separate audio channels – theoretically, our app could act as a "virtual interpreter" by joining the meeting audio and broadcasting a translated speech to an interpretation channel. This would be a clever but complex integration (would

require two Zoom connections: one to listen as an attendee, one to speak as the interpreter). - **Local Loopback for Output:** If we keep it simpler, the app might **play the translated audio locally** for the user (e.g. in their headphones) without injecting it into Zoom. That way, the user hears the translation of what others say. If the user also wants their speech translated for others, they could speak and let the app output the translated speech on a **virtual microphone** which is set as the input to Zoom. For example, the user speaks English -> app captures and translates to Russian -> app feeds the Russian audio into Zoom as if it were the user's microphone. This configuration can be achieved with virtual audio cable drivers and is technically feasible, though it takes some setup.

Considering **latency**, a 1–2 second delay is acceptable in conversation, but we have to manage turn-taking. The app should ideally mute the translated audio when the original speaker continues, to avoid talking over them. If the conversation is slow-paced (with pauses for responses), the slight delay is fine. In rapid back-and-forth, the overlapping might occur (similar to how UN interpreters sometimes have to cut in). This is a usability challenge but solvable with good UX cues (e.g. maybe a push-to-talk for translating your speech, etc.).

APIs for this use case: The same cloud APIs (Google, Azure, AWS) can be used in streaming mode here. Azure's Speech Translation API supports meetings with multiple speakers and can automatically handle language switching, outputting a chosen target language for all foreign speech ²³ ²⁴ – a perfect match for a bilingual meeting scenario. We could configure it to say “whatever language is spoken, translate it to English for me” or vice versa. This reduces complexity; the API would continuously output English audio for any non-English input. If both directions are needed, we might run two sessions (one translating Eng→Rus, another Rus→Eng). Another aspect is Zoom often has many jargon or names – we might use features like **custom glossaries** in translation (many APIs allow a glossary to ensure certain terms or names aren't mistranslated ²⁵). This could improve accuracy in professional meetings.

Competitors & demand: There is clear demand for this – language barriers in virtual meetings are common. Besides Zoom's own caption feature, Microsoft **Teams** and **Skype** have translation features. Skype Translator (since ~2017) can do real-time voice translation in about **10 spoken languages including Russian** ²⁶ , demonstrating the viability. Microsoft also rolled out live captions translation in Teams for some languages. Third-party solutions like **Wordly** provide an AI translator specifically for meetings – Wordly's service can join Zoom and provide **live audio translation in 30+ languages** as well as captions ²⁷ ²⁸ . In fact, Wordly allows each participant to choose a language and hear the meeting in that language on their device. This shows that our envisioned app is on the right track and technically doable – Wordly built it using presumably the same kind of pipeline (likely using cloud APIs under the hood). The competitive landscape here includes other startups and perhaps hardware solutions (e.g. custom translator devices for conferences). However, many of these either focus on captions or require enterprise setup. An easy-to-use app for consumers or small businesses to get translated audio in Zoom could stand out.

One should note privacy considerations: piping meeting audio to cloud services could raise confidentiality concerns. Some enterprise users might prefer on-premise models for that reason. But at least technically, using the cloud APIs yields the best performance. Perhaps offer an option to use a local mode (with open-source models) for sensitive meetings, acknowledging that quality might be lower.

Other Applications and Future Extensions

Beyond YouTube and Zoom, the same real-time translation capability can apply to **numerous scenarios**:

- **Live Streams and Webinars**: Any live video or audio stream (Twitch streams, live webinars, even live TV news) could be fed through the translator. For example, a user could input a livestream URL and get an overlay audio translation. This could be valuable for live international events or conferences. It's essentially the generalization of the YouTube live case. Many organizations run live events with translation – currently they hire interpreters or use captioning; an AI solution could scale to more languages at lower cost.
- **Other Language Pairs**: While we focused on English ↔ Russian, the system could extend to many languages. The underlying APIs support dozens; for instance, Zoom's caption service supports 35 languages ²⁹, Google Meet's latest update supports **69 languages** for translated captions (which amounts to 4,600+ possible language pairs) ³⁰. Our app could later include options for other common pairs (e.g. English-Spanish, English-Chinese, etc.) by simply changing API parameters. The design is largely language-agnostic as long as ASR and TTS support the languages. Russian/English is a good test case because both languages are well-supported and it covers a challenging script difference (Cyrillic to Latin).
- **Mobile/Portable Use**: Envision this technology on a smartphone or AR glasses. Someone could, say, use their phone in "conversation mode" to talk with a person in real-time (each person speaks their language and hears the translation of the other). Google Translate's mobile app already has a Conversation mode that does turn-by-turn voice translation for many languages. Our application could be adapted into a mobile app or a wearable that continuously translates in near real-time. The demand for travel translators or even multilingual meetings in-person is high. (Microsoft's example scenarios for Azure Speech include a multilingual meeting with delegates each using a device to talk and hear in their own language ³¹ ³².)
- **Enhancing with Video/Context**: For recorded videos, an interesting extension is to also **translate on-screen text or slides** (OCR + translate) and maybe even **re-sync the speaker's lip movements** via deepfake tech to better match the dubbed audio (this is very advanced, but research exists for "speech-to-speech translation with lip synchronization"). These are not required for MVP, but future versions could aim for a more immersive experience (especially for video content).
- **Voice Customization**: As a future improvement, one could incorporate voice cloning so that the translated speech **sounds more like the original speaker**. There are research projects (e.g. by Google and Microsoft) on preserving the speaker's vocal characteristics across languages. Microsoft's latest video translation tool even mentions a "personal voice" feature in limited access ³³. While not trivial to implement, this could be a competitive differentiator if achieved – the user hears the speaker as if they spoke in the target language themselves. Early versions might simply allow choosing a similar voice tone (male/female, etc.), but the tech trend is moving toward retaining speaker identity.
- **Transcription and Captioning**: Even if the main output is audio, the app can also display translated **subtitles** simultaneously. This is essentially already obtained as the intermediate text from the MT step. Providing the user an option to see the live transcript (original and/or translated) could be useful for clarity (or for the deaf/hard-of-hearing communities, to whom this technology also opens access in multiple languages). Many competitors like Zoom and Wordly do provide text captions – we can too, as a supplementary feature.

In summary, the core technology can be repurposed wherever real-time cross-language communication is needed. The market for this spans from business meetings, online education (imagine attending a foreign lecture with live translation), customer support centers (live translating calls or chats), to entertainment and personal use. As hardware and network latency improve (5G, etc.), real-time translation is becoming more seamless.

Market Demand and Competitors (Brief)

There is a **strong market demand** for real-time translation. Globalization of work and content has created many use cases: multinational companies want team meetings where everyone can speak their native tongue; content creators want larger audiences; travelers and language learners seek instant communication tools. This demand is evidenced by big tech companies investing in solutions: - **Zoom and Google Meet:** Both have introduced real-time translated captions. Zoom's translated caption add-on (launched 2022–2023) supports 30+ languages in text ²². Google Meet has gone even further, recently expanding to 69 languages for live translated captions ³⁰. These are text-based for now, but it's plausible they will add audio in the future (Google's experiment with translating speech in Google Glass in real time shows interest in live audio translation as well). - **Microsoft:** Long a leader in this space with **Skype Translator** and **Microsoft Translator**. Skype's real-time voice translator (using Microsoft's engine) has been around since mid-2010s, covering 10 spoken languages (including English, Russian, Chinese, etc.) ²⁶. Microsoft also offers **Translator Live** as a service where multiple users can join a conversation room and get cross-translations. And as noted, Azure Cognitive Services now offer ready-to-use speech translation APIs for developers. Microsoft's presence validates that the technology works and that there is user interest (Skype Translator was praised for its accuracy in many cases ³⁴). - **Google:** Google's Translator app on mobile can do live two-way voice translation for many languages (using its speech and text APIs under the hood). They also have research projects like **Translatotron** and **SeamlessM4T** (Meta's project) pushing the envelope on direct speech-to-speech translation without intermediate text. While not products yet, these indicate the trend and future potential. Google's YouTube has begun testing auto-dubbing for videos (a few channels now offer machine-generated dubbed audio tracks). - **Startups:** A number of startups and smaller companies are targeting this space. For example, **Wordly** (mentioned earlier) provides AI translation for meetings and events, with both subtitles and audio output in 30+ languages ²⁷ ²⁸. Another, **Interactio**, focuses on conference interpretation (historically with human interpreters, but they've been piloting AI). **Waverly Labs** sells earbuds that can translate conversations in near real-time for travelers. **Timekettle** is another that makes translator earbuds. These indicate a healthy competition to serve the need for communication across languages in real time. In video dubbing, startups like **Papercup**, **DeepDub**, and **Resemble AI** offer AI dubbing for recorded content (e.g. media companies translating their video libraries). They typically focus on quality over real-time, but it's part of the same market continuum – showing that many believe speech translation is a big opportunity.

Despite the competition, the field is not saturated – especially when it comes to combining **real-time + voice output + ease of use**. Many existing solutions either do captions only, or are enterprise-focused, or only handle a limited scenario. Our idea of a flexible app where you can input any audio source (file, video link, live call) and get instant voice translation is quite cutting-edge. The feasibility is supported by all the advancements discussed. Given acceptable latency and using APIs, development is achievable without inventing new algorithms from scratch.

Conclusion

In conclusion, a real-time English↔Russian translation application is **highly feasible** with current technology. By leveraging API services for speech recognition, translation, and speech synthesis, one can develop an app that translates audio with only a minor delay. The plan to start with offline audio file translation ("crawl"), then live translation in a local environment ("walk"), and finally a web-deployed service for any stream ("run") is prudent. Each step is technically attainable: the first is almost trivial with today's cloud APIs; the second has been demonstrated by cloud provider examples; the third is an integration

challenge but clearly within reach as shown by existing solutions like Wordly and ongoing R&D by big players. Market demand is confirmed by the flurry of activity in this space – users *want* to break language barriers in media and communication. By focusing on a robust technical implementation (perhaps using Google/Azure/AWS services initially) we can deliver a working product quickly. The competitors' presence further validates the idea, but also means we should aim to differentiate (perhaps via better user experience, more languages, or faster integration with platforms like YouTube/Zoom). Overall, **real-time speech translation** is not only possible – it's already here in early forms – and our proposed application can confidently be built to ride this wave of innovation. ⁸ ²⁸

Sources: The feasibility statements and examples above are supported by current documentation and industry reports, including cloud provider technical blogs (AWS, Azure) outlining speech-to-speech translation pipelines ⁸ ¹³, official updates from Zoom/Google on language support ²² ³⁰, and media articles on tools like Skype Translator and Wordly ²⁶ ²⁸, among others. All evidence points to a green light for proceeding with development of this real-time translation app.

¹ ² ⁴ ⁸ ⁹ ¹⁰ ¹⁶ ²⁰ ²⁵ Break through language barriers with Amazon Transcribe, Amazon Translate, and Amazon Polly | AWS Machine Learning Blog

<https://aws.amazon.com/blogs/machine-learning/break-through-language-barriers-with-amazon-transcribe-amazon-translate-and-amazon-polly/>

³ Introducing Whisper | OpenAI

<https://openai.com/index/whisper/>

⁵ ¹⁴ Best Azure Speech Translation Alternatives & Competitors

<https://sourceforge.net/software/product/Azure-Speech-Translation/alternatives>

⁶ ⁷ ¹¹ ¹² ¹³ ²³ ²⁴ ³¹ ³² ³³ Announcing Video Translation & Speech Translation API Enhancements | Microsoft Community Hub

<https://techcommunity.microsoft.com/blog/azure-ai-services-blog/announcing-video-translation--speech-translation-api-enhancements/4148007>

¹⁵ Top DeepL Alternatives in 2025

<https://slashdot.org/software/p/DeepL/alternatives>

¹⁷ Dubbing videos with Aloud to overcome the language barrier

<https://blog.google/technology/area-120/aloud/>

¹⁸ Use automatic dubbing - YouTube Help

<https://support.google.com/youtube/answer/15569972?hl=en>

¹⁹ YouTube has started rolling out Google Translate dubbed audio ...

https://www.reddit.com/r/youtube/comments/16i4btq/youtube_has_started_rolling_out_google_translate/

²¹ Solved: Zoom AI Companion real time translation

<https://community.zoom.com/t5/Zoom-AI-Companion/Zoom-AI-Companion-real-time-translation/m-p/157719>

²² ²⁹ Bridge the language gap and connect people with Zoom translated captions | Zoom

<https://www.zoom.com/en/blog/translated-captions/>

²⁶ Skype Can Now Translate Your Voice Calls Into 10 Different Languages in Real-Time

<https://futurism.com/skype-can-now-translate-your-voice-calls-into-10-different-languages-in-real-time>

27 28 Wordly Expands Zoom Translation Options with Mobile Support

<https://www.wordly.ai/blog/wordly-adds-zoom-translation-mobile>

30 Google Workspace Updates: Google Meet adds 52 new languages to translated captions and 11 more languages for closed captions

<http://workspaceupdates.googleblog.com/2024/06/more-languages-for-captions-and-translated-captions-in-google-meet.html>

34 Skype real time voice translator is so good (when working) - Reddit

https://www.reddit.com/r/skype/comments/1e6ukiz/skype_real_time_voice_translator_is_so_good_when/